

UNIVERSITY OF EUROPE FOR APPLIED SCIENCES

MSc Data Science Programme

Predicting Song Popularity from Audio Features: A Supervised Machine Learning Approach

Supervised Machine Learning — Assignment 1 Proposal

Submitted by: Nithin Chowdary Raya

Programme: MSc Data Science

University: University of Europe for Applied Sciences

Supervisor: Prof. Raja Hashim Ali

Submission Date: May 4th 2026

Contents

- 1 Project Title 2**
- 2 Dataset Information 2**
- 3 Target Variable 2**
 - 3.1 Primary Target — popularity (Regression) 3
- 4 Problem Statement 3**
- 5 Research Questions 4**
- 6 Proposed Methodology 4**
 - 6.1 Data Preprocessing 4
 - 6.2 Feature Engineering 5
 - 6.3 Modelling Pipeline 5
- 7 Machine Learning Models to be Used 6**
- 8 Evaluation Metrics 7**
- 9 Expected Figures and Tables 7**
 - 9.1 Research Question 1 — Feasibility of Popularity Prediction 7
 - 9.2 Research Question 2 — Feature Correlation with Popularity 8
 - 9.3 Research Question 3 — Algorithm Comparison 9
 - 9.4 Research Question 4 — Feature Engineering Impact 11
 - 9.5 Research Question 5 — Audio Features Across Popularity Levels 12
 - 9.6 Research Question 6 — Effect of Training Data Size 13
 - 9.7 Research Question 7 — Evaluation Metrics Comparison 14
- References 15**
- A Dataset Feature Summary 16**

1. Project Title

Predicting Song Popularity from Audio Features Using Supervised Machine Learning

2. Dataset Information

This analysis uses a large-scale audio features dataset containing approximately one million songs with 14 attributes per record, including audio features and metadata. Key dataset attributes are summarised in Table 1.

Table 1: Dataset Summary

Attribute	Details
Dataset Name	Almost Million Songs Dataset 2025 (16 Features)
Source	Kaggle — anantsinghal786/almost-million-songs-dataset-2025-16-features
Dataset Link	https://www.kaggle.com/datasets/anantsinghal786/almost-million-songs-dataset-2025-16-features
Total Records	~1,000,000 individual song records
Total Features	14 columns covering audio features, genre, and artist metadata
Format	Comma-Separated Values (.csv)
Year Published	2025

The dataset encompasses 14 attributes per song record, capturing audio characteristics (*tempo*, *energy*, *danceability*, *valence*, *acousticness*, *instrumentalness*, *liveness*, *speechiness*, *loudness*, *key*, *mode*, *time_signature*), metadata (*genres*, *track_artists*), and the target variable **popularity** (0–100). This richness makes it well-suited to regression tasks for predicting song popularity.

3. Target Variable

3.1 Primary Target — popularity (Regression)

Each row in the dataset represents a single song. The numeric column `popularity` is a Spotify-derived score ranging from 0 to 100, reflecting the relative current popularity of the track based on total play count and recency of streams. With values distributed across the full range but skewed towards lower scores, this continuous target makes the task a regression problem, directly shaping research questions around feature correlation, model comparison, and evaluation metrics.

Table 2: Task Type Definition

Task Type	Target Column	Justification
Regression (Primary)	<code>popularity</code>	Continuous numeric score (0–100) representing song popularity

4. Problem Statement

The music streaming business produces massive amounts of information, and the popularity of songs on. Spotify and other platforms is a key business indicator to artists, labels, and playlist editors. Although this is important, the connection between the quantifiable audio properties of a song is not easy. and its popularity rating is still not fully comprehended. The majority of previous research is concerned with genre. classification or mood detection, and the actual prediction of popularity based on audio features are left to the problem solver. underexplored at scale. This project explores the idea that machine learning regression models can be able to predict accurately. popularity of songs based on 14 audio and metadata features, with a dataset of about one. million songs published in 2025. This task is technically challenging due to three interrelated challenges. non-trivial.

1. The popularity distribution is extremely skewed - most of the songs have low ratings. which has a long tail of popular tracks - which requires careful consideration beyond mere. mean error reporting.
2. Audio characteristics (energy, danceability, and valence) may have non-linear interactions, need ensemble models and feature engineering to express these relationships.
3. Identifying which audio features are genuinely predictive versus spuriously correlated requires rigorous feature importance analysis and cross-validation.

Insights from this project have practical applications in music recommendation systems, playlist curation, and supporting independent artists in understanding the sonic attributes

associated with greater listener reach.

5. Research Questions

Seven targeted research questions guide the investigation:

RQ1:

Can supervised regression models accurately predict song popularity from audio features?

RQ2:

Which audio features have the strongest correlation with song popularity?

RQ3:

How do different regression algorithms compare in predicting song popularity?

RQ4:

How does feature engineering (e.g., polynomial features) affect model prediction accuracy for song popularity?

RQ5:

How do audio features vary across different popularity levels?

RQ6:

Does increasing training data size improve model generalisation when predicting song popularity?

RQ7:

Which evaluation metrics (MAE, RMSE, R^2) best characterise model performance for song popularity prediction?

6. Proposed Methodology

6.1 Data Preprocessing

- Missing values in `track_artists` and `genres` will be imputed with a sentinel value 'unknown' to retain all song records.
- Label encoding will be applied to the `genres` column: `genre_count` (number of genres per track) and `top_genre` (most frequent genre label) will be derived as numeric features.

- The `track_artists` column will be dropped from modelling features after encoding, as raw artist names are too high-cardinality for direct use.
- Feature scaling using `StandardScaler` will be applied for distance-based and linear models; tree-based models will use unscaled data.
- An 80/20 stratified train-test split (stratified on popularity quartile) will ensure consistent target distribution across both partitions.

6.2 Feature Engineering

Table 3: Engineered Features

Engineered Feature	Type	Derivation / Purpose
<code>genre_count</code>	Integer	Number of genres listed per song; captures genre diversity
<code>energy_x_danceability</code>	Float	Interaction term; captures songs that are both high-energy and danceable
<code>loudness_abs</code>	Float	Absolute loudness value; removes sign for cleaner distribution
<code>energy_minus_acousticness</code>	Float	Contrast between electronic and acoustic production style
<code>dance_x_valence</code>	Float	Happy and danceable songs; associated with mainstream appeal
<code>tempo_sq</code>	Float	Polynomial tempo feature to capture non-linear tempo effects
<code>speech_x_live</code>	Float	Interaction of speechiness and liveness; identifies live spoken-word content

6.3 Modelling Pipeline

Data will flow through the following sequential stages:

1. **Data collection and loading:** Downloading the dataset from Kaggle and loading into a Pandas DataFrame.
2. **Preprocessing:** Encoding genres, imputing missing values, deriving engineered features, and applying train-test split.

3. **Exploratory Data Analysis:** Generating distribution plots, correlation heatmaps, and popularity-level audio feature comparisons.
4. **Model training and hyperparameter tuning:** Training all regression models using cross-validated RandomizedSearchCV.
5. **Evaluation and comparison:** Computing all metrics, generating feature importance plots, and producing expected figures and tables.

7. Machine Learning Models to be Used

Table 4: Regression Models for popularity Prediction

Model	Configuration / Rationale
Linear Regression	Parametric baseline; establishes minimum performance floor for popularity prediction
Ridge Regression	Regularised linear model controlling for multicollinearity among correlated audio features
Decision Tree Regressor	Rule-based learner; provides human-interpretable decision paths through feature space
Random Forest Regressor	Bagging ensemble; primary source of feature importance rankings; reduces overfitting
Gradient Boosting Regressor	Sequential boosting; strong predictive accuracy for tabular regression tasks
K-Nearest Neighbours (k=10)	Instance-based non-parametric model; good for local similarity patterns in audio space

8. Evaluation Metrics

Table 5: Evaluation Metrics

Metric	Formula	Rationale
MAE	$\frac{1}{n} \sum \hat{y}_i - y_i $	Interpretable: average prediction error in popularity points
RMSE	$\sqrt{\frac{1}{n} \sum (\hat{y}_i - y_i)^2}$	Penalises large errors more than MAE; sensitive to outliers
R ² Score	$1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$	Proportion of variance in popularity explained by the model (0–1)
Median Absolute Error	Median $ \hat{y}_i - y_i $	Robust to outliers; appropriate given skewed popularity distribution

9. Expected Figures and Tables

This section presents one representative figure and one data table for each of the seven research questions. All outputs use consistent styling and are labelled according to the naming convention defined in this proposal.

9.1 Research Question 1 — Feasibility of Popularity Prediction




Figure 1: Distribution of Song Popularity Scores (0–100)




Figure 2: Actual vs Predicted Popularity: Naive Baseline vs Linear Regression vs Random Forest

Table 6: Table 1.1 — Model Performance Summary (Test Set)

Model	MAE	RMSE	R ²
Naive Baseline (mean)	~22.0	~27.5	0.000
Linear Regression	~20.5	~26.0	~0.09
Random Forest Regressor	~17.2	~22.8	~0.24

Indicative values; exact figures will be reported after model training on the full dataset.

9.2 Research Question 2 — Feature Correlation with Popularity




Figure 3: Pearson Correlation of Audio Features with Popularity (Sorted)




Figure 4: Random Forest Feature Importance for Popularity Prediction

Table 7: Table 2.1 — Feature Ranking by Method

Feature	Pearson Rank	MI Rank	RF Rank
danceability	1	3	2
energy	3	2	3
loudness	4	4	4
tempo	5	5	5
acousticness	6	6	6

Indicative rankings; exact values computed from the dataset during analysis.

9.3 Research Question 3 — Algorithm Comparison




Figure 5: Model Performance Comparison: MAE, RMSE, R² (All Models)




Figure 6: Actual vs Predicted Popularity Scatter Plots (All Models)

Table 8: Table 3.1 — Algorithm Comparison

Model	MAE	RMSE	R ²	Time (s)
Linear Regression	~20.5	~26.0	~0.09	<1
Ridge Regression	~20.5	~26.0	~0.09	<1
Decision Tree	~18.0	~23.5	~0.18	<5
Random Forest	~17.2	~22.8	~0.24	~30
Gradient Boosting	~17.0	~22.5	~0.25	~60
KNN (k=10)	~19.0	~24.5	~0.14	~10

Indicative values; exact figures will be reported after model training.

9.4 Research Question 4 — Feature Engineering Impact




Figure 7: MAE Before vs After Feature Engineering (Ridge and Random Forest)




Figure 8: Performance with Polynomial Features (Degree 2, Ridge Regression)

Table 9: Table 4.1 — Feature Engineering Impact

Configuration	MAE	RMSE	R ²
Ridge — Original Features	~20.5	~26.0	~0.09
Ridge — Engineered Features	~19.8	~25.2	~0.12
Ridge — Polynomial (deg 2)	~19.0	~24.5	~0.15
RF — Original Features	~17.2	~22.8	~0.24
RF — Engineered Features	~16.8	~22.2	~0.26

Indicative values; exact figures will be reported after model training.

9.5 Research Question 5 — Audio Features Across Popularity Levels




Figure 9: Box Plots: Audio Features Across Popularity Quartiles




Figure 10: Heatmap of Mean Audio Feature Values by Popularity Band

Table 10: Table 5.1 — Mean Feature Values by Popularity Band

Feature	Low (0–25)	Med-Low (25–50)	Med-High (50–75)	High (75–100)
danceability	~0.52	~0.58	~0.64	~0.70
energy	~0.55	~0.58	~0.61	~0.65
acousticness	~0.35	~0.28	~0.22	~0.16
valence	~0.46	~0.50	~0.54	~0.58
tempo (BPM)	~118	~121	~122	~124

Indicative values; exact figures computed from dataset during EDA.

9.6 Research Question 6 — Effect of Training Data Size




Figure 11: Learning Curves: MAE and R² vs Training Data Size




Figure 12: Marginal MAE Gain per Training Size Increase (Random Forest)

Table 11: Table 6.1 — Performance vs Training Size

Training Size	Ridge MAE	Ridge R ²	RF MAE	RF R ²
500	~23.0	~0.01	~22.5	~0.03
5,000	~21.5	~0.06	~19.0	~0.15
20,000	~20.8	~0.08	~17.8	~0.21
100,000	~20.5	~0.09	~17.2	~0.24
150,000	~20.4	~0.09	~17.0	~0.25

Indicative values; exact figures will be reported after model training.

9.7 Research Question 7 — Evaluation Metrics Comparison




Figure 13: All Evaluation Metrics Across All Regression Models




Figure 14: Radar Chart: Top 3 Models Across Key Metrics

Table 12: Table 7.1 — Comprehensive Metric Comparison

Model	MAE	RMSE	R ²	MedAE	MAPE%	MaxErr
Linear Regression	~20.5	~26.0	~0.09	~17.0	~28%	~98
Random Forest	~17.2	~22.8	~0.24	~14.5	~24%	~96
Gradient Boosting	~17.0	~22.5	~0.25	~14.2	~23%	~95

Indicative values; exact figures will be reported after model training.

References

- Almost Million Songs Dataset 2025 (16 Features). *Kaggle*. Retrieved from <https://www.kaggle.com/datasets/anantsinghal786/almost-million-songs-dataset-2025-16-feat>
- Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (2nd ed.). Springer.
- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5–32.
- Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785–794.

A. Dataset Feature Summary

Table 13: Dataset Feature Summary

Column Name	Data Type	Description
genres	String	List of genre labels associated with the track
track_artists	String	Name(s) of the performing artist(s)
tempo	Float	Estimated beats per minute of the track
energy	Float (0–1)	Perceptual measure of intensity and activity
key	Integer (0–11)	Musical key of the track (0 = C, 1 = C#, etc.)
popularity	Integer (0–100)	Spotify popularity score — the regression target variable
mode	Integer (0/1)	Modality: 1 = major, 0 = minor
time_signature	Integer	Estimated time signature (e.g., 3 = 3/4, 4 = 4/4)
speechiness	Float (0–1)	Presence of spoken words; higher values indicate more speech-like content
danceability	Float (0–1)	Suitability for dancing based on tempo, rhythm, and beat regularity
valence	Float (0–1)	Musical positiveness; higher = more happy and euphoric
acousticness	Float (0–1)	Confidence measure of whether the track is acoustic
liveness	Float (0–1)	Detects presence of an audience; higher = more likely live recording
instrumentalness	Float (0–1)	Predicts whether a track has no vocals; higher = more instrumental

Dataset: <https://www.kaggle.com/datasets/anantsinghal786/almost-million-songs-datas>